


Anyone can start an AI agent. I build the part that says no.

Building enterprise systems since 1998, and hands-on with generative AI since 2022. Avanade, Microsoft and Lloyd's of London, then co-founder and CEO of SynapseDx. AI is moving fast, and startups learn fastest close to real problems. So alongside SynapseDx, I take on a small number of fractional enterprise AI and architecture engagements.



- The agent asks to book a meeting. That's within its rules, so it goes through.
- It asks to pay a £40,000 invoice. That isn't, so it's stopped and a person decides.

Where I'm useful

Running AI in production, not just launching it.

Agents fail quietly. I set up the alarms that notice when one goes silent, a record of every action it takes, and a clear path for the people who have to act when something goes wrong.

Build and deploy a team can trust.

Every change checked before it goes in, released the same way to every environment, and new environments set up from a description rather than by hand. It holds whether a person or an AI wrote the code.

Guardrails for agents and for engineers.

Agents get only the access they need, every action is checked against their rules, and if the check can't run the answer is no. Tests have to be seen failing before anyone trusts them to pass.

Architecture that holds up.

One home for every fact, decisions written down where people will find them, and review boards that help teams ship rather than queue.

What I've built, and what it taught me

Four decisions from building SynapseDx. You'll face the same ones.



Every agent action goes through one gate.

Nothing an AI agent does reaches a business system without being checked against its rules and written down. If the gate can't check, the answer is no.



Write the decision down, then let a check enforce it.

Architecture decisions live where people actually read them. An automatic check stops anyone, person or agent, from quietly breaking one.



The engineering system comes first. Then the AI writes code.

Every piece of work starts as a written design and a plan, which a second model reviews before any code exists. Tests have to fail before they're trusted to pass, and nothing merges until every check is green. The agents type faster. The architecture, the rules and the sign-off stay human.



Silence is not a pass.

No error, no alert and no output can all mean something broke. A check only counts once it has been seen catching the real failure. It's the rule I'd bring to any team running AI.

Skills

Enterprise architecture

AI agents and guardrails

Data architecture

Leading architecture teams

Running AI in production

Build and deploy pipelines

Architecture governance

Large language models

Microsoft Azure

Databricks

Power BI and analytics

Engineering systems for AI coding agents

Identity and access

Python and TypeScript

Postgres

Google Cloud

Pre-sales and deal shaping

Board and fundraising

Agile delivery

Models I build with, by the value each brings

60%

Claude builds: design, code and agents

Gemini challenges: a voice in the multi-model council

25%

GPT and Codex review: a second opinion on every plan and change

Grok and DeepSeek cross-check

10%

5%

Several side by side, so one model's blind spot never becomes the system's.

The last fourteen years

SynapseDx

Co-Founder & CEO, lead platform architect. Aug 2024 to now, full-time from Jul 2026

Governing AI agents

Designed the gateway every agent action passes through, so each one is checked against its rules and recorded. Set up how code gets designed, checked and released by a team that works with AI agents. Run the seed raise, the board and go-to-market.

Microsoft

Director, Cloud Solution Architecture (Data & AI). Dec 2021 to Mar 2023

Data and AI at scale

Led three teams of 45 cloud solution architects in data and AI, behind an \$80m-a-year business in industrials, manufacturing, energy and pharma.

How to be early and wrong about AI agents in production

Apollo Product Engineering, September 2026

Beyond the AI hype

The esynergy Podcast, August 2026

Lloyd's of London

Lead Enterprise Data Architect. Sep 2023 to Jun 2026. Joined Endava in Mar 2023, placed at Lloyd's through Endava, then direct from Apr 2024

Technology in a regulated market

Reshaped the architecture review board and design authority that decide what gets built in the world's specialist insurance market. Evaluated new technology, including AI, set the plans for APIs, disaster recovery and test environments, and ran the Microsoft and Databricks relationships.

Avanade

Director, Digital Analytics, Innovation & Industry. Apr 2012 to Feb 2020

Machine learning before it was called AI

Built early machine learning prototypes on Azure and spoke on predictive analytics in healthcare to 5,000 people. Delivered the first Power BI implementation in North America, and led data programmes and deals up to \$3.1m.

Broker bot, meet underwriter bot

Early 2023, to a big tech audience. AI agents as broker and underwriter, with a human underwriter working across both

Predictive analytics in healthcare

Avanade Tech Summit

Got an agent nobody will sign off?
Let's talk.

maria@synapsedx.ai

bines.ai

